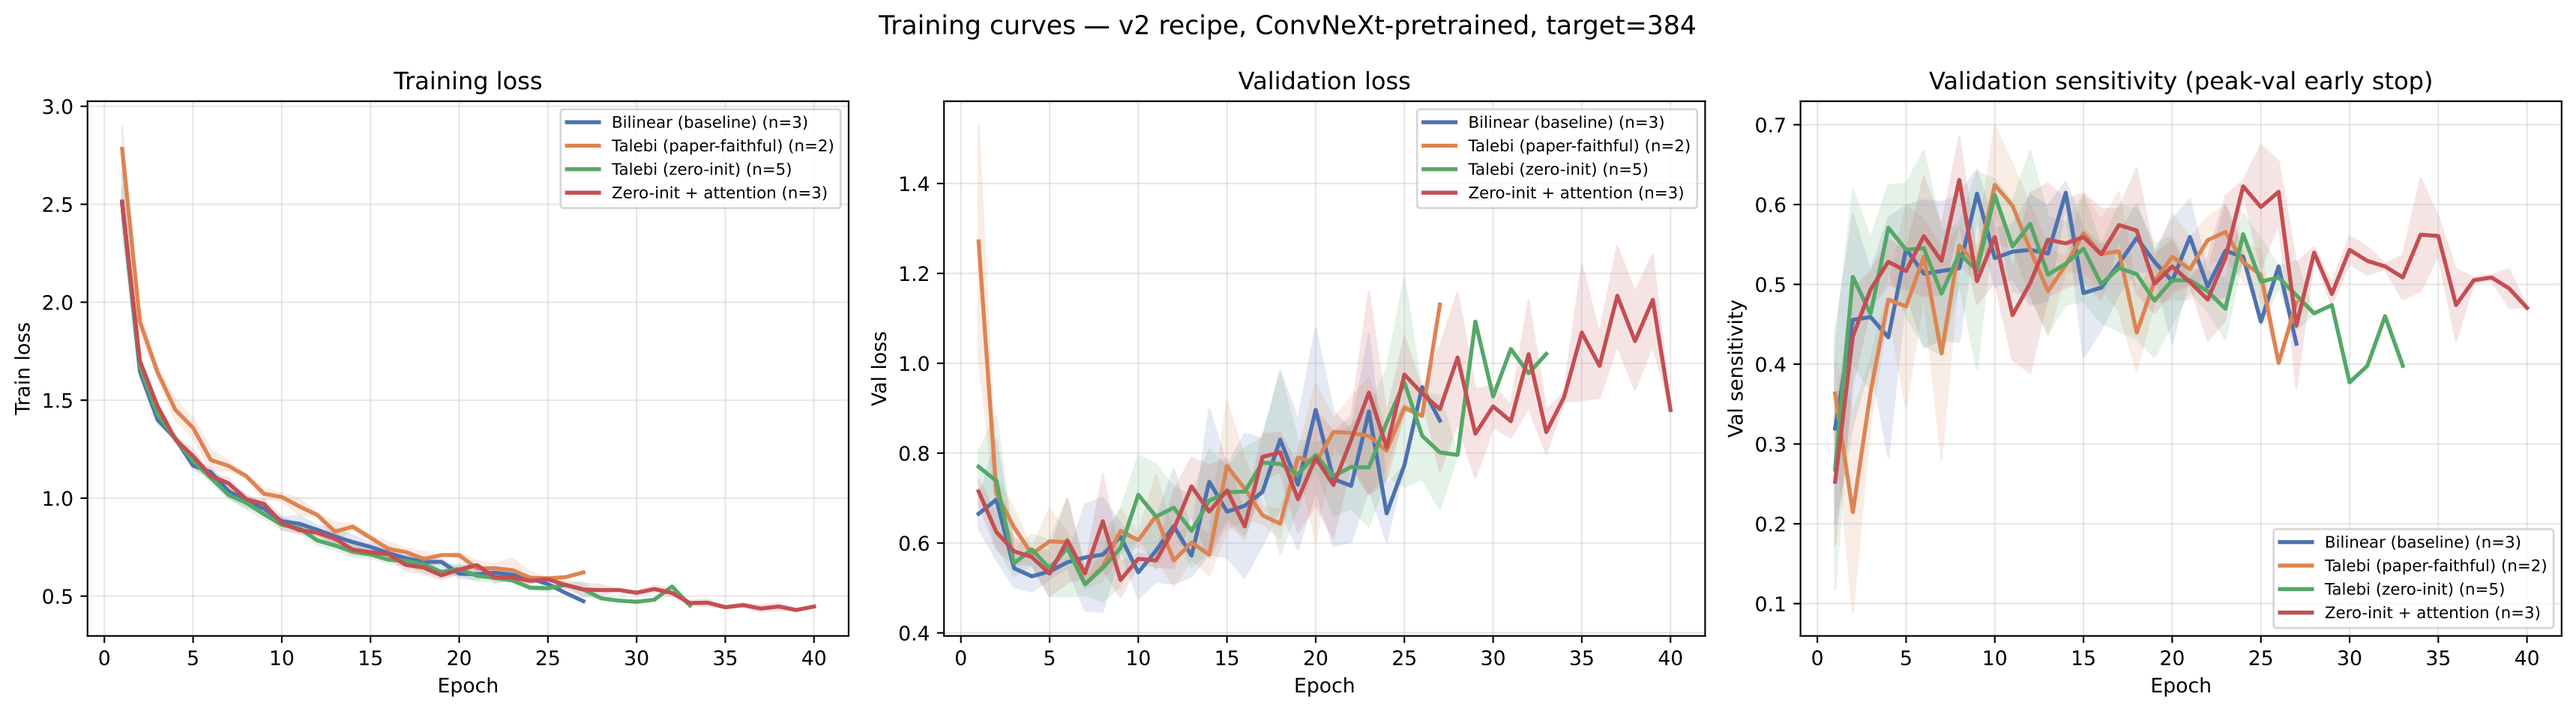
Training curves — v2 recipe, ConvNeXt-pretrained, target=384
Training loss
Train loss
Epoch
Validation loss
Val loss
Epoch
Validation sensitivity (peak-val early stop)
Val sensitivity
Epoch
Bilinear (baseline) (n=3)
Talebi (paper-faithful) (n=2)
Talebi (zero-init) (n=5)
Zero-init + attention (n=3)